

A nicer numpy

Dima Kogan

February 29, 2020

What is this about?

Two libraries to make working in numpy nicer.

- ▶ These are public tools, available for some years
- ▶ Installable from Debian and related distros.
- ▶ Python2 and Python3 both supported

`numpysane` (<https://github.com/dkogan/numpysane>)

- ▶ Nicer array manipulation
- ▶ Nicer basic linear algebra routines
- ▶ Broadcasting support

`gnuplotlib` (<https://github.com/dkogan/gnuplotlib>)

- ▶ Plotting

What's wrong with numpy?

- ▶ Some very core functionality is mysterious and unintuitive.
- ▶ Things work as expected *only* with 2-dimensional arrays, no more and no less.

Areas addressed by `numpysane`:

- ▶ Nicer array manipulation
- ▶ Nicer basic linear algebra routines
- ▶ Broadcasting support

Matrix concatenation

Very easy example: stick two identical 2D arrays together

► The docs say to use `hstack()` or `vstack()` or `dstack()`

```
>>> import numpy as np
>>> arr32 = np.arange(3*2).reshape(3,2)
>>> print(arr32)
[[0 1]
 [2 3]
 [4 5]]

>>> print(arr32.shape)
(3, 2)
```

Matrix concatenation

What do we expect `hstack(arr32, arr32)` to do?

```
[[0 1 0 1]
 [2 3 2 3]
 [4 5 4 5]]
```

or

```
[[0 1]
 [2 3]
 [4 5]
 [0 1]
 [2 3]
 [4 5]]
```

Matrix concatenation

This was a trick question. Here's what it does:

```
>>> np.hstack(arr32,arr32)
```

```
Traceback (most recent call last):
```

```
  File "<stdin>", line 1, in <module>
```

```
TypeError: hstack() takes 1 positional argument ...
```

```
... but 2 were given
```

Apparently `hstack()` wants an iterable of the arguments, instead of the arguments themselves

Matrix concatenation

Fine. Here's what it does if you feed it what it wants:

```
>>> print(np.hstack((arr32, arr32)))  
[[0  1  0  1]  
 [2  3  2  3]  
 [4  5  4  5]]
```

Whew. That makes sense! Looks "horizontal".

Matrix concatenation

What if I don't feed it strictly 2D matrices?

```
>>> arr132 = np.arange(3*2).reshape(1,3,2)
>>> print(arr132)
[[[0 1]
  [2 3]
  [4 5]]]

>>> print(arr132.shape)
(1, 3, 2)
```


Matrix concatenation

Same question as before: what do we expect
`hstack((arr132,arr132))` to do?

```
[[[0 1 0 1]
   [2 3 2 3]
   [4 5 4 5]]]
```

or

```
[[[0 1]
   [2 3]
   [4 5]
   [0 1]
   [2 3]
   [4 5]]]
```

or something else?

Matrix concatenation

Here's what it does:

```
>>> print(np.hstack((arr132,arr132)))  
[[[0 1]  
  [2 3]  
  [4 5]  
  [0 1]  
  [2 3]  
  [4 5]]]
```

```
>>> np.hstack((arr132,arr132)).shape  
(1, 6, 2)
```

Whoa. That is *not* horizontal at all! I would have expected a result with shape (1,3,4)

Matrix concatenation

What if I give it 1-dimensional arrays?

```
>>> arr3 = np.arange(3)
>>> arr13 = np.arange(3).reshape(1,3)
>>> print(arr3)
[0 1 2]

>>> arr3.shape
(3,)

>>> print(arr13)
[[0 1 2]]

>>> arr13.shape
(1, 3)
```

Matrix concatenation

```
>>> np.hstack((arr3, arr3)).shape  
(6,)
```

```
>>> np.hstack((arr13, arr13)).shape  
(1, 6)
```

```
>>> np.hstack((arr13, arr3)).shape  
ValueError: all the input arrays must have ...  
... same number of dimensions
```

Looks like the easy 1D cases work reasonably well, but it can't

```
hstack([[0 1 2]], [0 1 2])
```

I claim it *should* be able to do this, and will justify it later. And other functions such as `vstack()` do accept arrays with uneven input dimensions.

Matrix concatenation

So what's wrong?

- ▶ **numpy is inconsistent about which is the most significant dimension in an array**
- ▶ There's an arbitrary design choice that must be made: if I stack N arrays of shape (A, B, C) into a new array, do I get an array of shape (N, A, B, C) or (A, B, C, N) ?
- ▶ Most of numpy makes the *first* choice, but some of it (concatenation functions most notably) makes the second choice

Dimensionality example

Example:

- ▶ Let's say I have a 1-dimensional array containing simultaneous temperature measurements at different locations:

```
>>> print(temperatures)
[ t_place0 t_place1 t_place2 ... ]
```

```
>>> print(temperatures.shape)
(Nlocations,)
```

To access temperatures at different locations in this array, I look along the first dimension (`axis = 0`). This is a 1-dimensional array, so that's also the last dimension (`axis = -1`).

Dimensionality example

Let's now say I have measured all the temperatures multiple times throughout the day, and I make a joint array using the *first* convention above

```
>>> print(temperatures)
[[ t_time0place0 t_time0place1 t_time0place2 ... ]
 [ t_time1place0 t_time1place1 t_time1place2 ... ]
 ...]
```

```
>>> print(temperatures.shape)
(Ntimes, Nlocations)
```

- ▶ I want to use the same axis index to access different locations
- ▶ If I count dimensions from the end, I can do that: `axis = -1` refers to the different locations in both cases
- ▶ But if I count dimensions from the front, this doesn't work: locations *were* at `axis = 0` but now they're at `axis = 1`

Dimensionality example

If I used the *second* convention above to produce an array of shape (Nlocations, Ntimes) then the locations are at axis = 0 in both cases. *But* the arrays would be printed differently:

- ▶ the location axis would be printed horizontally in the 1D array, and vertically in the 2D array

So:

- ▶ numpy printing effectively counts dimensions from the back
- ▶ `hstack()` counts dimensions from the front. It *should be* concatenating along axis = -1 (the last one)
- ▶ But instead it concatenates along axis = 1 (the 2nd one) with special-case rules about what to do with 1-dimensional arrays that don't have a 2nd axis
- ▶ For 2D arrays the 2nd axis *is* the last axis, but for 3D arrays it isn't

Matrix concatenation with numpysane

There are two functions, both stolen from the PDL project.

- ▶ `glue()` concatenates any N arrays along the given axis
- ▶ `cat()` concatenates N arrays along a new outer dimension

These both add leading length-1 dimensions to the input as needed: "something" is logically equivalent to "1 of something". This is one of the *broadcasting* rules I'll get to in a bit

Matrix concatenation with numpysane

`nps.glue()` works as expected:

```
>>> import numpysane as nps
```

```
>>> nps.glue(arr32, arr32, axis=-1).shape  
(3, 4)
```

```
>>> nps.glue(arr32, arr32, axis=-2).shape  
(6, 2)
```

```
>>> nps.glue(arr132, arr132, axis=-1).shape  
(1, 3, 4)
```

```
>>> nps.glue(arr13, arr3, axis=-1).shape  
(1, 6)
```

```
>>> nps.glue(arr13, arr3, axis=-2).shape  
(2, 3)
```

Matrix concatenation with numpysane

`nps.cat()` works as expected too. It always adds a new leading dimension

```
>>> nps.cat(arr32, arr32).shape  
(2, 3, 2)
```

```
>>> nps.cat(arr132, arr32).shape  
(2, 1, 3, 2)
```

Matrix multiplication

The funny business extends to other core areas of numpy. For instance multiplying matrices is non-trivial

- ▶ Up until numpy 1.10.0 `np.dot()` was the function for that, and it is surprising in all sorts of ways
- ▶ In 1.10.0 we got `np.matmul`, which is *much* better, but even then it has strange corners. Trying to compute an outer product:

```
>>> a = np.arange(5).reshape(5,1)
```

```
>>> b = np.arange(3)
```

```
>>> np.matmul(a,b)
```

```
ValueError: matmul: Input operand 1 has a mismatch in  
its core dimension 0, with gufunc signature  
(n?,k),(k,m?)->(n?,m?) (size 3 is different from 1)
```

Matrix multiplication with numypsane

numypsane provides its own `matmult()` routine that does what one expects:

```
>>> nps.matmult(a,b).shape
```

```
(5, 3)
```

There're many more functions in numypsane in this area. Everything's documented, and I'd like to move on to...

Broadcasting

What is broadcasting?

- ▶ *Broadcasting* is a generic way to vectorize functions
- ▶ A broadcasting-aware function has a *prototype*: it knows the dimensionality of its inputs and of its outputs
- ▶ When calling a broadcasting-aware function, any extra dimensions in the input are automatically used for vectorization

Broadcasting: an example

This is best described with an example: a broadcasting-aware inner product. An inner product (also known as a dot product) is a function that

- ▶ takes in two identically-sized 1-dimensional arrays
- ▶ outputs a scalar

If one calls a broadcasting-aware inner product with two arrays of shape (2,3,4) as input, it would

- ▶ compute 6 inner products of length-4 each
- ▶ report the output in an array of shape (2,3)

Because `nps.inner()` knows the dimensionality of its inputs and of its outputs, it can figure out how to parse the input arrays

Broadcasting: an example

```
>>> a234 = np.arange(2*3*4).reshape(2,3,4)
```

```
>>> print(a234)
```

```
[[[ 0  1  2  3]
   [ 4  5  6  7]
   [ 8  9 10 11]]
```

```
 [[12 13 14 15]
   [16 17 18 19]
   [20 21 22 23]]]
```

```
>>> print(np.inner(a234,a234))
```

```
[[ 14 126 366]
 [ 734 1230 1854]]
```

The values in the output are `inner([0,1,2,3], [0,1,2,3])` and `inner([4,5,6,7], [4,5,6,7])` and so on.

What about the stock broadcasting support?

What stock broadcasting support?

- ▶ It's sparse and incomplete
- ▶ Little end-user awareness that it exists

Broadcasting rules

In short:

- ▶ Line up the shapes of the inputs to their *trailing* dimensions
- ▶ Match the trailing dimensions with the expected shapes of the inputs
- ▶ Any leading dimensions left over are used for vectorization
- ▶ The extra leading dimensions must be compatible across all the inputs. This means that each leading dimension must either
 - ▶ equal 1
 - ▶ be missing (thus assumed to equal 1)
 - ▶ equal to some positive integer >1 , consistent across all arguments
- ▶ The leading dimensions of the inputs determine the shape of the output

Broadcasting: an example

Let's write a broadcasting-aware inner product.

```
import numpy as np
@np.broadcast_define( (('n',), ('n',)), () )
def inner(a,b):
    # We could use numpy for this: return a.dot(b)
    sum = 0.
    for i in range(len(a)): sum += a[i]*b[i]
    return sum
```

Broadcasting: an example

- ▶ We wrote a function `inner(a,b)` that computes *one* inner product. It knows nothing about vectorization
- ▶ And it can assume that `a` and `b` are 1-dimensional arrays of the same length
- ▶ Then we used the `nps.broadcast_define()` decorator to add broadcasting awareness to this function
- ▶ The decorator is told that there are 2 inputs, that have dimensions `('n',)` each. It's the same `n` for both of them, so ...
- ▶ the internal `nps.broadcast_define()` machinery will make sure that the *last* dimension of both inputs matches
- ▶ If we call `inner()` with more than 1-dimensional input, we'll get multiple inner products computed, and an array of output returned

Broadcasting: an example

So we can give it two arrays, and get inner products of each corresponding row:

```
>>> a234 = np.arange(2*3*4).reshape(2,3,4)
```

```
>>> print(inner(a234,a234).shape)  
(2,3)
```

Or we can compute the inner product of each row of one array to some arbitrary vector

```
>>> a234 = np.arange(2*3*4).reshape(2,3,4)
```

```
>>> a4    = np.arange(4)
```

```
>>> print(inner(a234,a4).shape)  
(2,3)
```

Broadcasting: more involved example

Let's say we have a function with input prototype

► ((3,), ('n',3), ('n',), ('m',))

given inputs of shape

(1,5, 3)

(2,1, 8,3)

(8)

(5, 9)

The broadcasting logic will set $n = 8$ and $m = 9$.

The call will return an output array of shape $(2,5, \dots)$, where \dots is the shape of each output slice.

Broadcasting: summary

- ▶ This is a very powerful technique. The `nps.broadcast_define()` decorator is written in Python and wraps Python code. With lots of iterations this is *slow*.
- ▶ A much faster analogue exists in C: `nps.numpysane_pywrap()`. The iteration code and the code for the inner function are all in C, so this is fast. Please see the documentation for more detail.
- ▶ A stock nump broadcasting-in-C API exists:
<https://docs.scipy.org/doc/numpy-1.13.0/reference/c-api.generalized-ufuncs.html>
I found this after implementing my own, and have not tried it.

Plotting: gnuplotlib

Let's switch gears, and talk about plotting.

- ▶ As with the numpy core, there's a dominant choice here: matplotlib
- ▶ I'm not aware of any *major* issues: if it's not pissing you off right now, there probably isn't a lot of reason to switch to my library

However, numpy ...

- ▶ is python-specific
- ▶ is slow
- ▶ has a weird API
- ▶ is missing useful interactivity

Plotting: `gnuplotlib`

`gnuplotlib`: a plotting library for `numpy`

- ▶ Uses `gnuplot` as the plotting backend, so
 - ▶ The plots look and interact like `gnuplot` plots have for decades
 - ▶ It's fast
 - ▶ Lots of features and backends available
- ▶ Has a (claimed) reasonable API
- ▶ A direct port of `PDL::Graphics::Gnuplot`

Plotting: gnuplotlib design choices

- ▶ One `plot()` function does everything
 - ▶ Can still build up the plot components programmatically with python
- ▶ `gnuplotlib` is a thin shim: strings are passed to `gnuplot` verbatim (like in `feedgnuplot`)
 - ▶ So we get a powerful library and a friendly learning curve


Introductory tutorial is available at

<https://github.com/dkogan/gnuplotlib/blob/master/guide/guide.org>

Plotting: gnuplotlib: a *very* brief tutorial

To plot something, just call plot:

```
import numpy          as np
import numpysane      as nps
import gnuplotlib     as gp
th = np.linspace(-2.*np.pi, 2.*np.pi, 100)
gp.plot(np.sin(th))
```



Plotting: gnuplotlib: a *very* brief tutorial


```
import numpy          as np
import numpysane      as nps
import gnuplotlib     as gp
th = np.linspace(-2.*np.pi, 2.*np.pi, 100)
gp.plot(np.sin(th))
```

- ▶ We're plotting in 2D, so default is `tuplesize=2` arrays
- ▶ We gave it just 1 array, so integers 0,1,2,... were used for the x

Plotting: gnuplotlib: a *very* brief tutorial

- We can pass in 2 arrays to make a scatter x-y plot:

```
th = np.linspace(-np.pi, np.pi, 100)  
gp.plot(np.cos(th), np.sin(th), square = True)
```



Plotting: gnuplotlib: a *very* brief tutorial

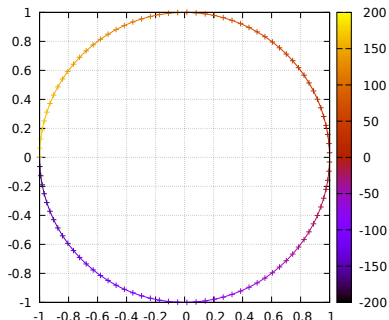
```
th = np.linspace(-np.pi, np.pi, 100)
gp.plot(np.cos(th), np.sin(th), square = True)
```

- ▶ We passed in *two* arrays
- ▶ We also passed in `square = True`. This is a *plot option* to autoscale the x and y axes evenly. Otherwise the circle will look like an ellipse

Plotting: gnuplotlib: a *very* brief tutorial

- It's possible to have more values per point. For instance:

```
th = np.linspace(-np.pi, np.pi, 100)
gp.plot(np.cos(th), np.sin(th),
        # The angle (in degrees) is shown as the color
        th * 180./np.pi,
        tuplesize = 3,
        _with      = 'linespoints palette',
        square     = True)
```



Plotting: gnuplotlib: a *very* brief tutorial

```
th = np.linspace(-np.pi, np.pi, 100)
gp.plot(np.cos(th), np.sin(th),
        # The angle (in degrees) is shown as the color
        th * 180./np.pi,
        tuplesize = 3,
        _with      = 'linespoints palette',
        square     = True)
```

- ▶ The style `linespoints palette` is given to gnuplot directly. `gnuplotlib` doesn't know what that means
- ▶ `tuplesize=3` tells `gnuplotlib` that there are 3 values per point. Because of `palette`, these will be interpreted as `x,y,color`
- ▶ The `gnuplot` documentation talks in detail about what kind of input each style expects

Plotting: gnuplotlib: a *very* brief tutorial

An explicit invocation of `plot()` looks like this:

```
plot( curve, curve, ..., plot_options )
```


where each curve is a tuple:

```
curve = (array, array, ..., curve_options)
```

- ▶ *plot options* apply to the whole plot, and are given as keyword args to `plot()`
- ▶ *curve options* apply to each separate curve (dataset); given in a `dict()` in the end of each curve tuple. Or defaults given in the `plot()` kwargs
- ▶ If we have one dataset, we can inline the tuples, like we did above

Plotting: gnuplotlib: a *very* brief tutorial

```
th = np.linspace(-2.*np.pi, 2.*np.pi, 100)
gp.plot( ( th, np.sin(th), ),
        ( th, np.cos(th), ),
        ( th, th, dict(_with = 'points ps 1') ),
        _with = 'lines',
        xlabel = "Angle (rad)",
        title  = "2 with lines and 1 with points")
```



Plotting: gnuplotlib: a *very* brief tutorial


```
th = np.linspace(-2.*np.pi, 2.*np.pi, 100)
gp.plot( ( th, np.sin(th), ),
        ( th, np.cos(th), ),
        ( th, th, dict(_with = 'points ps 1') ),
        _with = 'lines',
        xlabel = "Angle (rad)",
        title  = "2 with lines and 1 with points")
```

- ▶ We passed in 3 tuples, one for each dataset
- ▶ We passed in the `xlabel` plot option to label the x axis
- ▶ We passed in the `title` plot option to title the plot
- ▶ We passed in the default `with` curve option: `lines`
- ▶ 2/3 datasets don't set their own `with`, so they use `lines`
- ▶ 1/3 plots with `points ps 1` instead. `gnuplotlib` doesn't know what that is, but `gnuplot` knows that `ps` is a synonym for `pointsize`

Plotting: gnuplotlib: a *very* brief tutorial

- Broadcasting is fully supported:

```
th = np.linspace(-2.*np.pi, 2.*np.pi, 100)
gp.plot( th,
         nps.cat(np.sin(th),
                 np.cos(th)),
         legend = np.array( ("sin", "cos"), ) )
```



Plotting: gnuplotlib: a *very* brief tutorial

```
th = np.linspace(-2.*np.pi, 2.*np.pi, 100)
gp.plot( th,
         nps.cat(np.sin(th),
                 np.cos(th)),
         legend = np.array( ("sin", "cos"), ) )
```

- ▶ I plotted two datasets, but didn't use tuples
- ▶ Using default `tuplesize=2`, and gave it two arrays:
 - ▶ First array has the expected shape of `(100,)`
 - ▶ Second array has the shape `(2,100)`
- ▶ This thus broadcasts: I get two plots: `sin(th)` vs `th` and `cos(th)` vs `th`
- ▶ curve options broadcast too: I have it two different legend options, and `gnuplotlib` knows to use each one for the two datasets


Plotting: gnuplotlib: a *very* brief tutorial

► What does this do?

```
th      = np.linspace(0, 6*np.pi, 200)
z       = np.linspace(0, 5,          200)
size    = 0.5 + np.abs(np.cos(th))
color   = np.sin(2*th)

gp.plot3d( np.cos(th) * nps.transpose(np.array((1,-1))),
           np.sin(th) * nps.transpose(np.array((1,-1))),
           z,
           size,
           color,
           tuplesize = 5,
           _with = 'points ps variable pt 7 palette',
           squarexy = True)
```

Plotting: gnuplotlib: a *very* brief tutorial



Plotting: gnuplotlib

More or less that's it. Lots of examples in the guide:

- ▶ <https://github.com/dkogan/gnuplotlib/blob/master/guide/guide.org>

The API docs are on the main page:

- ▶ <https://github.com/dkogan/gnuplotlib>

Thanks for listening!

The documentation and sources and links to this talk:

- ▶ <https://github.com/dkogan/numypsane>
- ▶ <https://github.com/dkogan/gnuplotlib>

Or you can

```
apt install python3-numypsane python3-gnuplotlib
```